

# Reinforcement Learning – Part 2



# Outline

- **Value-based RL**
- TD Learning
- SARSA
- Q-learning
- Deep RL



# Reinforcement Learning

- **Model-based RL**

- Learn the reward function and transition probabilities
- Use planning (e.g., Value Iteration)
- Extract policy

- **Value-based RL**

- Learn the value function directly (e.g.,  $Q^*$ )
- Extract policy



- **Policy-based RL**

- Learn the policy directly

# Value-based RL

- “Most popular” RL methods
- Include **many** methods
  - E.g., Monte Carlo methods, TD methods, etc.

# Outline

- Value-based RL
- **TD Learning**
- SARSA
- Q-learning
- Deep RL



# TD Learning

- Temporal-difference (TD) learning is a model-free RL approach
  - Learn by **bootstrapping** from the current estimate (just like Dynamic Programming)
  - Learn by **sampling the environment** (just like Monte Carlo methods)

# TD Learning

- What is the difference between model-based RL and model-free RL?
  - **Model-based RL:** learns the transition probabilities and reward function + planning
  - **Model-free RL:** opposite (e.g., learns the  $V^*(s)$  or  $Q^*(s, a)$  directly)

# TD Learning

- TD(0) – Estimating  $V^\pi(s)$ 
  - For every new  $(s_t, r_{t+1}, s_{t+1})$
  - Update

$$V(s_t) \leftarrow V(s_t) + \alpha[r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

---

*TD Error*



# TD Learning

- *TD Error*

$$\delta_t = r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$$

Where

$$V(s_t)$$

is the estimated value of  $s_t$ ,

$$r_{t+1} + \gamma V(s_{t+1})$$

is the better estimate

# TD Learning

- We also want to use TD prediction for the **control problem** (i.e., finding an optimal policy)
  - **On-policy method** – estimate function (e.g.,  $Q^\pi(s, a)$ ) for the current behavior policy  $\pi$
  - **Off-policy method** – estimate function (e.g.,  $Q^{\pi'}(s, a)$ ) for a different policy  $\pi'$  than the behavior policy  $\pi$

# Outline

- Value-based RL
- TD Learning
- **SARSA**
- Q-learning
- Deep RL



# SARSA

- SARSA is an on-policy TD Control

- For every new  $(s_t, a_t, r_{t+1}, s_{t+1}, a_{t+1})$

- Update

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

---

*TD Error*

# SARSA

- We start with some estimate  $Q$
- Initialize current state  $s$  and choose some action  $a$  (e.g., using  $\epsilon$ -greedy)
- Loop for each step:
  - Take action  $a$  and observe next state  $s'$  and reward  $r$
  - choose some action  $a'$  (e.g., using  $\epsilon$ -greedy)
  - Update  $Q$  estimate according to
$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$$
  - $s \leftarrow s'$
  - $a \leftarrow a'$

# Outline

- Value-based RL
- TD Learning
- SARSA
- **Q-learning**
- Deep RL



# Q-learning

- Q-learning is an off-policy TD Control
  - For every new  $(s_t, a_t, r_{t+1}, s_{t+1})$
  - Update

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \underbrace{\left[ r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t) \right]}_{TD \text{ Error}}$$

# Q-learning

- We start with some estimate  $Q$
- Initialize current state  $s$
- Loop for each step:
  - Choose some action  $a$  (e.g., using  $\epsilon$ -greedy)
  - Take action  $a$  and observe next state  $s'$  and reward  $r$
  - Update  $Q$  estimate according to
$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$
- $s \leftarrow s'$



# Value-based RL

- Does Q-learning work?
- Theorem: As long as **every state-action pair** is visited infinitely often, Q-learning converges to  $Q^*$  w.p.1.

Let us  
revisit  
this



# Exploration vs. Exploitation

- **How can we visit every state action pair infinitely often?**
  - In practice, “infinitely often” means a “large number of times”

# Exploration vs. Exploitation

- The agent needs to **try all actions in all states many times**
- But this means that the agent will keep **doing sub-optimal actions for a long time!**

# Exploration vs. Exploitation


On the other hand...

# Exploration vs. Exploitation

- We want the agent to start **acting “reasonably” as soon as possible**
  - In practice, this means **using knowledge already available**
  - The agent needs to stop “trying” and start “doing”

# Exploration vs. Exploitation

- The agent needs to balance:
  - **Exploration:** trying new actions or actions that have not been selected very often
  - **Exploitation:** using learned knowledge to select actions



**Exploration vs. Exploitation  
tradeoff**

# Exploration vs. Exploitation

- Heuristic for exploration vs. exploitation
  - $\epsilon$ -greedy
    - Agent selects a random action with probability  $\epsilon$  (exploration)
    - Agent selects the greedy action (i.e., action with highest Q-value) with probability  $1 - \epsilon$  (exploitation)



$\epsilon$  may decay with time

# Example

- We have the following MDP:

- $S = \{1, 2, 3\}$


- $A = \{left, right\}$

- ~~■  $P(s'|s, a = left) = ?$~~

- ~~■  $P(s'|s, a = right) = ?$~~

- ~~■  $R(s, a) = ?$~~

- $\gamma = 0.9$





# Example

- We start with some estimate  $Q$

$$Q = \begin{bmatrix} 1 & 1 \\ 1 & 1 \\ 1 & 1 \end{bmatrix}$$

- Initialize current state  $s$

$$s \rightarrow 1$$

# Example

- First iteration (current state  $s \rightarrow 1$ )
  - Choose some action  $a$ 
    - $a \rightarrow left$
  - Take action  $a$  and observe next state  $s'$  and reward  $r$ 
    - $s' \rightarrow 1$
    - $r \rightarrow 0$
  - Update  $Q$  estimate according to
    - $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
    - $Q(1, left) \leftarrow 1 + 0.3[0 + 0.9 \times 1 - 1]$
    - $Q(1, left) \leftarrow 0.97$

# Example

- Updated  $Q$

$$Q = \begin{bmatrix} 0.97 & 1 \\ 1 & 1 \\ 1 & 1 \end{bmatrix}$$

# Example

- Second iteration (current state  $s \rightarrow 1$ )
  - Choose some action  $a$ 
    - $a \rightarrow right$
  - Take action  $a$  and observe next state  $s'$  and reward  $r$ 
    - $s' \rightarrow 2$
    - $r \rightarrow 0$
  - Update  $Q$  estimate according to
    - $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
    - $Q(1, right) \leftarrow 1 + 0.3[0 + 0.9 \times 1 - 1]$
    - $Q(1, right) \leftarrow 0.97$

# Example

- Updated  $Q$

$$Q = \begin{bmatrix} 0.97 & 0.97 \\ 1 & 1 \\ 1 & 1 \end{bmatrix}$$

# Example

- Third iteration (current state  $s \rightarrow 2$ )
  - Choose some action  $a$ 
    - $a \rightarrow left$
  - Take action  $a$  and observe next state  $s'$  and reward  $r$ 
    - $s' \rightarrow 1$
    - $r \rightarrow 0$
  - Update  $Q$  estimate according to
    - $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$
    - $Q(2, left) \leftarrow 1 + 0.3[0 + 0.9 \times 0.97 - 1]$
    - $Q(2, left) \leftarrow 0.9619$

# Example

- Updated  $Q$

$$Q = \begin{bmatrix} 0.97 & 0.97 \\ 0.9619 & 1 \\ 1 & 1 \end{bmatrix}$$

# Example

- Updated  $Q$  (after many iterations)

$$Q = \begin{bmatrix} 6.86 & 7.99 \\ 7.22 & 8.91 \\ 9.2 & 10 \end{bmatrix}$$

- Recall the  $Q^*$  with Value Iteration

$$Q^* = \begin{bmatrix} 6.89 & 7.66 \\ 7.08 & 8.73 \\ 9.07 & 9.95 \end{bmatrix}$$




# Example

- After computing  $Q$ , we can extract the policy as follows:

$$\pi^*(s) \in \operatorname{argmax}_{a \in A} Q(s, a)$$

$$\pi^* = \begin{bmatrix} \textit{right} \\ \textit{right} \\ \textit{right} \end{bmatrix}$$

$$\pi^* = \begin{bmatrix} 0 & 1 \\ 0 & 1 \\ 0 & 1 \end{bmatrix}$$



# Example

```
import numpy as np
np.set_printoptions(precision=2, suppress=True)

# States
S = ['1', '2', '3']

# Actions
A = ['L', 'R']

# Transition probabilities

L = np.array([[1.0, 0.0, 0.0],
              [0.8, 0.2, 0.0],
              [0.0, 0.8, 0.2]])

R = np.array([[0.2, 0.8, 0.0],
              [0.0, 0.2, 0.8],
              [0.0, 0.0, 1.0]])

P = [L, R]

# Reward function

R = np.array([[0.0, 0.0],
              [0.0, 0.0],
              [1.0, 1.0]])

gamma = 0.9
```

# Example

```
def egreedy(Q, state, eps):  
    p = np.random.random()  
  
    if p < eps:  
        action = np.random.choice(num_actions)  
    else:  
        action = np.argmax(Q[state,:])  
  
    return action
```

# Example

```
STEPS = 1000000
num_actions = len(A)
num_states = len(S)
ALPHA = 0.3

# Initialize Q-values
Q = np.ones((num_states, num_actions))

# Initialize current state
state = 0

for t in range(STEPS):

    # choose action
    action = egreedy(Q, state, 0.05)

    # choose next state
    next_state = np.random.choice(num_states, p=P[action][state, :])

    # obtain reward
    reward = R[state, action]

    # Update Q
    Q[state, action] = Q[state, action] + ALPHA*(reward + gamma*max(Q[next_state, :]) - Q[state, action])

    state = next_state

print(Q)
```

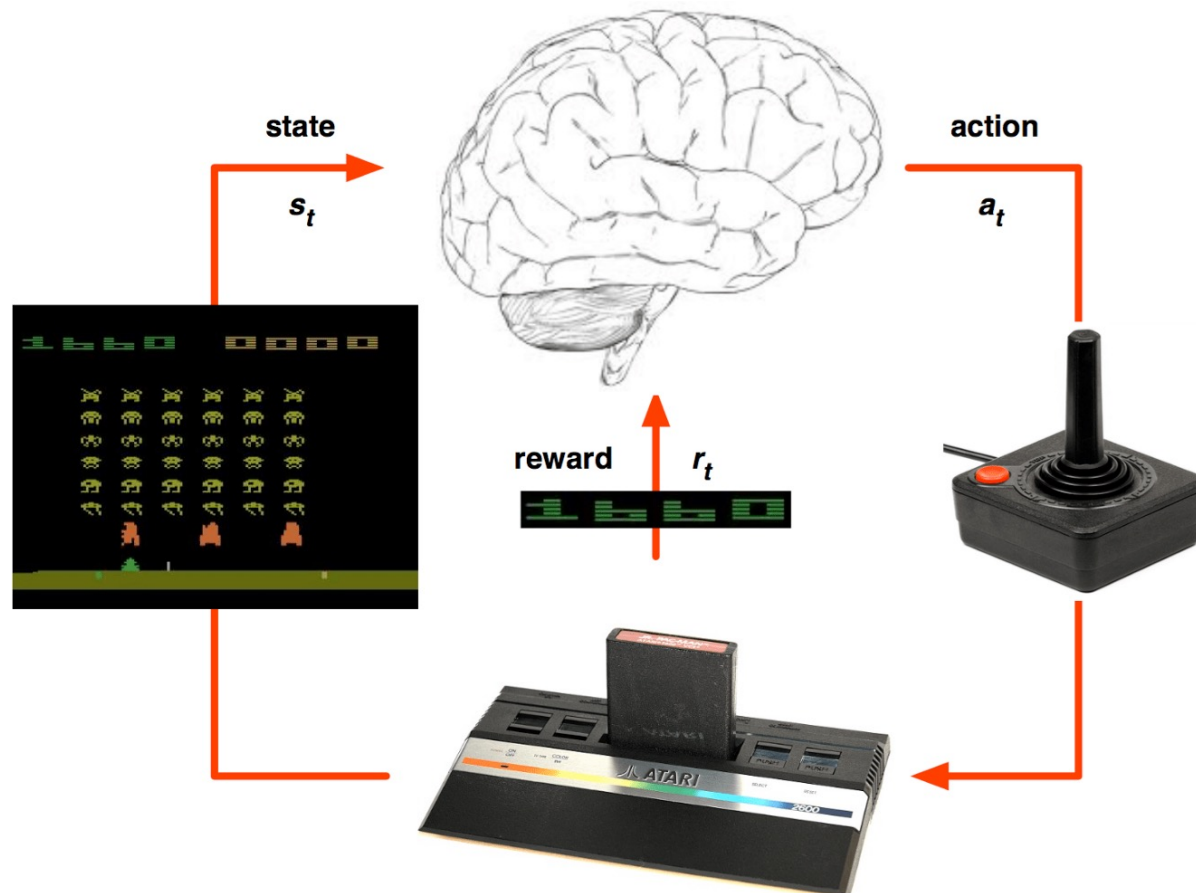
# Outline

- Value-based RL
- TD Learning
- SARSA
- Q-learning
- **Deep RL**



# Deep Reinforcement Learning


Google DeepMind's Deep Q-network (DQN)



# Deep Reinforcement Learning

- Used in domain with **large state space**
- We can **no longer represent  $Q^*$  exactly**
- We must resort to some **form of approximation** (e.g., neural network)
- Function approximation **does not retain its convergence guarantees**

# Deep Reinforcement Learning



State = vector of features



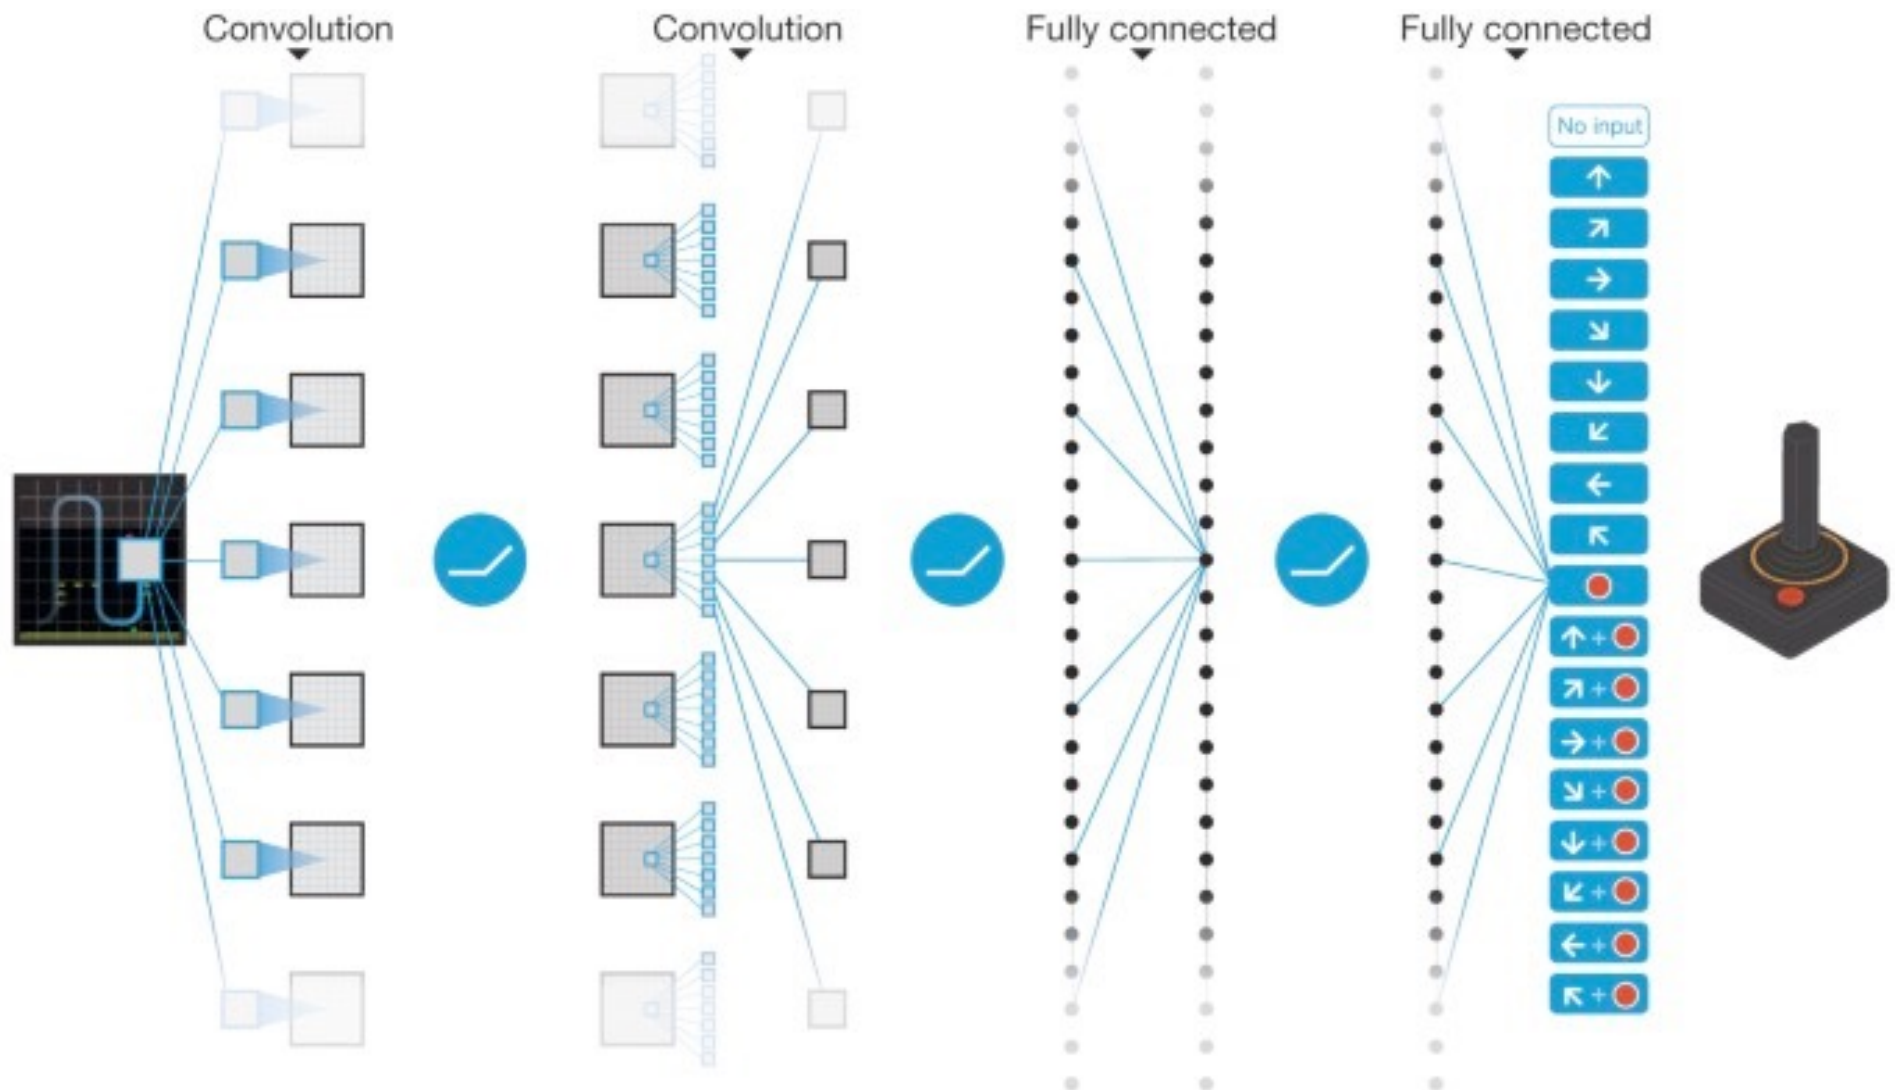
# Deep Q-networks


- Introduction to Deep Reinforcement Learning



<https://www.youtube.com/watch?v=wrBUkpiRvCA>

# Deep Reinforcement Learning





# DQN in Action



# Thank You



[sardinha@inf.puc-rio.br](mailto:sardinha@inf.puc-rio.br)